

Getting started with Kubernetes

In this lab you are going to learn the basics of Kubernetes. The lab uses kubectl together with a minikube cluster and is very easy to go through as it uses an already existing docker image.

If you want to do a more extensive getting started you can follow this [tutorial](#).

Prerequisites

Before you can execute the commands you need to have setup your environment on either [native windows 10](#) or using a [vagrant VM](#) with Ubuntu.

Getting started

First thing you need to do is start your minikube cluster if it is not yet running. Open `boxstarter shell` or `Powershell` on Windows 10 or `bash shell` on Ubuntu. To start your minikube cluster issue the command `minikube start`. You should see a similar result as below:

```
Starting local kubernetes v1.12.4 cluster...
Starting VM...

Everything looks great. Please enjoy minikube!
```

Testing if cluster is live

Test if you can connect to the `minikube` cluster from your command-line using `kubectl`. It is your local CLI command center. You can issue commands to the cluster either by `kubectl` or directly calling REST APIs exposed by the master. Issue the following command to get some info about your cluster.

```
kubectl cluster-info
```

The result of the command should be similar as below:

```
Kubernetes master is running at https://10.1.127.124:8443
kubeDNS is running at https://10.1.127.124:8443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.
```

How many nodes are part of your cluster?

If you want to know how many nodes (bare metal machines / VMs) are part of your cluster you can run the command `kubectl get nodes`. As we are using `minikube` this should have the following result.

```
$> kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
minikube	Ready	master	3d	v1.13.2

What pods are available by default?

The first concept of Kubernetes are **Pods**. A pod is *a group of one or more containers*. Pods are **the smallest deployable units** that can be created and managed in Kubernetes. Most of Kubernetes components also runs on pods. Issue the command `kubectl get pods` to get all available pods or search on a namespace.

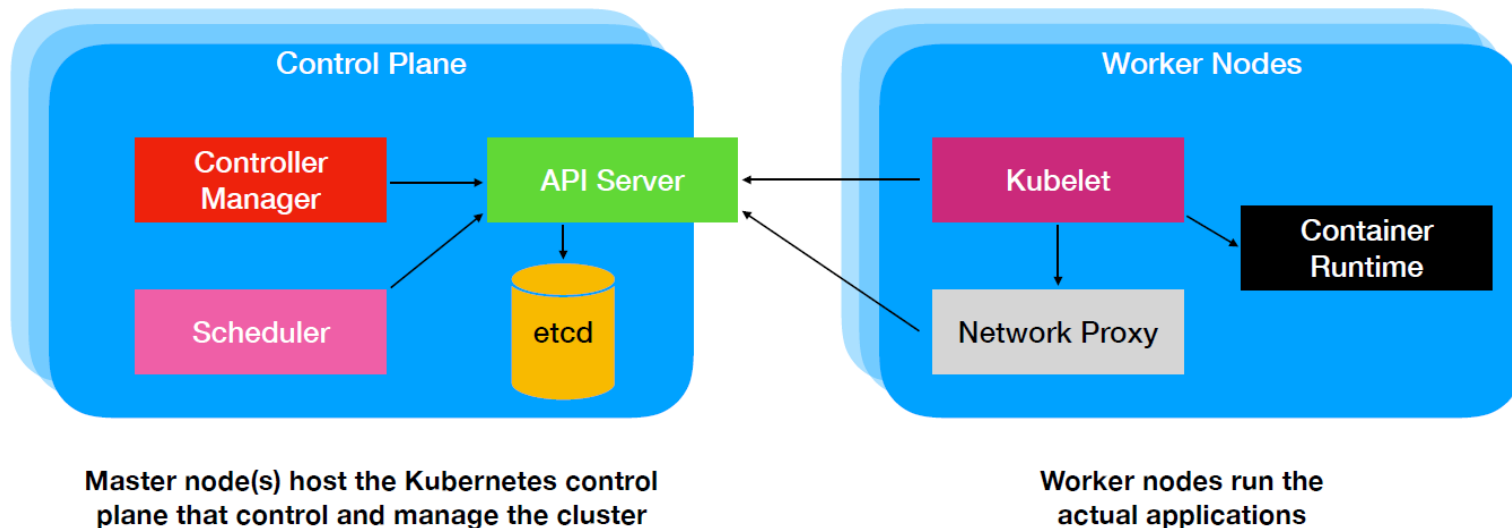
The second concept, **namespaces**, are intended for use in environments with many users spread across multiple teams, or projects. For clusters with a few to tens of users, you should not need to create or think about namespaces at all.

All pods by *kubernetes* and *minikube* are part of the `kube-system` namespace. You can see them by issuing the same command but with the addition of `-n kube-system`.

```
$> kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-576cbf47c7-qzmz4	1/1	Running	1	25m
coredns-576cbf47c7-swsrf	1/1	Running	1	25m
etcd-minikube	1/1	Running	0	3m
kube-addon-manager-minikube	1/1	Running	1	24m
kube-apiserver-minikube	1/1	Running	0	3m
kube-controller-manager-minikube	1/1	Running	0	3m
kube-proxy-kpbs9	1/1	Running	0	3m
kube-scheduler-minikube	1/1	Running	1	24m
kubernetes-dashboard-5bff5f8fb8-rvwbh	1/1	Running	1	24m
storage-provisioner	1/1	Running	1	24m

Can you spot the above pods in the below Kubernetes architecture?



Control Plane

Multiple components that can run on a single master node or be split across multiple (master) nodes and replicated to ensure high availability:

- **API Server:** communication center for developers, sysadmin and other Kubernetes components
- **Scheduler:** assigns a worker node to each deployable component
- **Controller Manager:** performs cluster-level functions (replication, keeping track of worker nodes, handling nodes failures...)
- **etcd:** reliable distributed data store where the cluster configuration is persisted

Worker Node

Machines that run containerized applications. It runs, monitors and provides services to applications via components:

- **Docker, rkt, or another container runtime:** runs the containers
- **Kubelet:** talks to API server and manages containers on its node
- **Network Proxy:** load balance network traffic between application components

Connect minikube cluster to local docker environment

Before you can deploy docker images on the Kubernetes container runtime you need to configure minikube so that it uses your local docker repository. This can be done by either issuing the command

- Powershell: `minikube docker-env | Invoke-Expression`
- Bash: `eval $(minikube docker-env)`

If the command is successful it returns an empty result.

Create a deployment on Kubernetes cluster

A *Deployment* controller provides declarative updates for [Pods](#) and [ReplicaSets](#) (a replica set (3rd concept of Kubernetes) ensures how many replica of pod should be running). You describe a *desired state* in a Deployment object, and the Deployment controller changes the actual state to the desired state at a controlled rate. You can define Deployments to create new ReplicaSets, or to remove existing Deployments and adopt all their resources with new Deployments.

In this lab we are going to install Nginx (a reverse proxy) as a deployment. Issue the following command in `powershell` or `bash`.

```
kubectl create deployment hello-nginx --image=nginx
```

This command looks in your local docker repository if the image exists, if not it is downloaded using docker, and deploys it as a pod on your cluster. You will get similar (final) output is shown below.

```
deployment.extensions "hello-nginx" created
```

To view your new deployment issue the command `kubectl get deployments`.

```
$> kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
hello-nginx	1	1	1	1	1m

View the pods the deployment has started by issuing the command `kubectl get pods`.

```
$> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
hello-nginx-6476b4b4-dnsrp	1/1	Running	0	21m

To understand what is happening in the background there is a handy command for this. Issue the command `kubectl get events` to get a similar result as below.

LAST SEEN	FIRST SEEN	COUNT	NAME	KIND	SUBOBJECT
	TYPE	REASON	SOURCE	MESSAGE	
2m	2m	1	hello-nginx-84d8d5d8f5-b8wkh.15779d3700ff4a2c	Pod	
	Normal	Scheduled	default-scheduler	Successfully assigned	
	default/hello-nginx-84d8d5d8f5-b8wkh to minikube				
2m	2m	1	hello-nginx-84d8d5d8f5-b8wkh.15779d372bdd96e0	Pod	
	spec.containers{nginx}	Normal	Pulling	kubelet, minikube	pulling image "nginx"
1m	1m	1	hello-nginx-84d8d5d8f5-b8wkh.15779d4909d2e2b4	Pod	
	spec.containers{nginx}	Normal	Pulled	kubelet, minikube	Successfully pulled
	image "nginx"				
1m	1m	1	hello-nginx-84d8d5d8f5-b8wkh.15779d49109067e8	Pod	
	spec.containers{nginx}	Normal	Created	kubelet, minikube	Created container
1m	1m	1	hello-nginx-84d8d5d8f5-b8wkh.15779d4918c345ac	Pod	
	spec.containers{nginx}	Normal	Started	kubelet, minikube	Started container
2m	2m	1	hello-nginx-84d8d5d8f5.15779d36fd7a9c58	ReplicaSet	
	Normal	SuccessfulCreate	replicaset-controller	Created pod: hello-nginx-	
	84d8d5d8f5-b8wkh				
2m	2m	1	hello-nginx.15779d36fa8c2c3c	Deployment	

Create a service for the deployment

Last step for getting the docker image started is by exposing the deployment as a service. A **service** is an abstraction which defines a logical set of `Pods` and a policy by which to access them - sometimes called a micro-service. The set of `Pods` targeted by a `Service` is (usually) determined by a [Label Selector](#).

1. Expose the Pod to the public internet using the `kubectl expose` command:

```
kubectl expose deployment hello-nginx --type=LoadBalancer --port=80
```

The `--type=LoadBalancer` flag indicates that you want to expose your Service outside of the cluster.

2. View the Service you just created:

```
kubectl get services
```

Output:

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-nginx	LoadBalancer	10.98.44.97	<pending>	80:30529/TCP	22s
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	3d

On cloud providers that support load balancers, an external IP address would be provisioned to access the Service. On Minikube, the `LoadBalancer` type makes the Service accessible through the `minikube service` command.

3. Finally, retrieve the URL of the deployed Nginx service by issuing the following command.

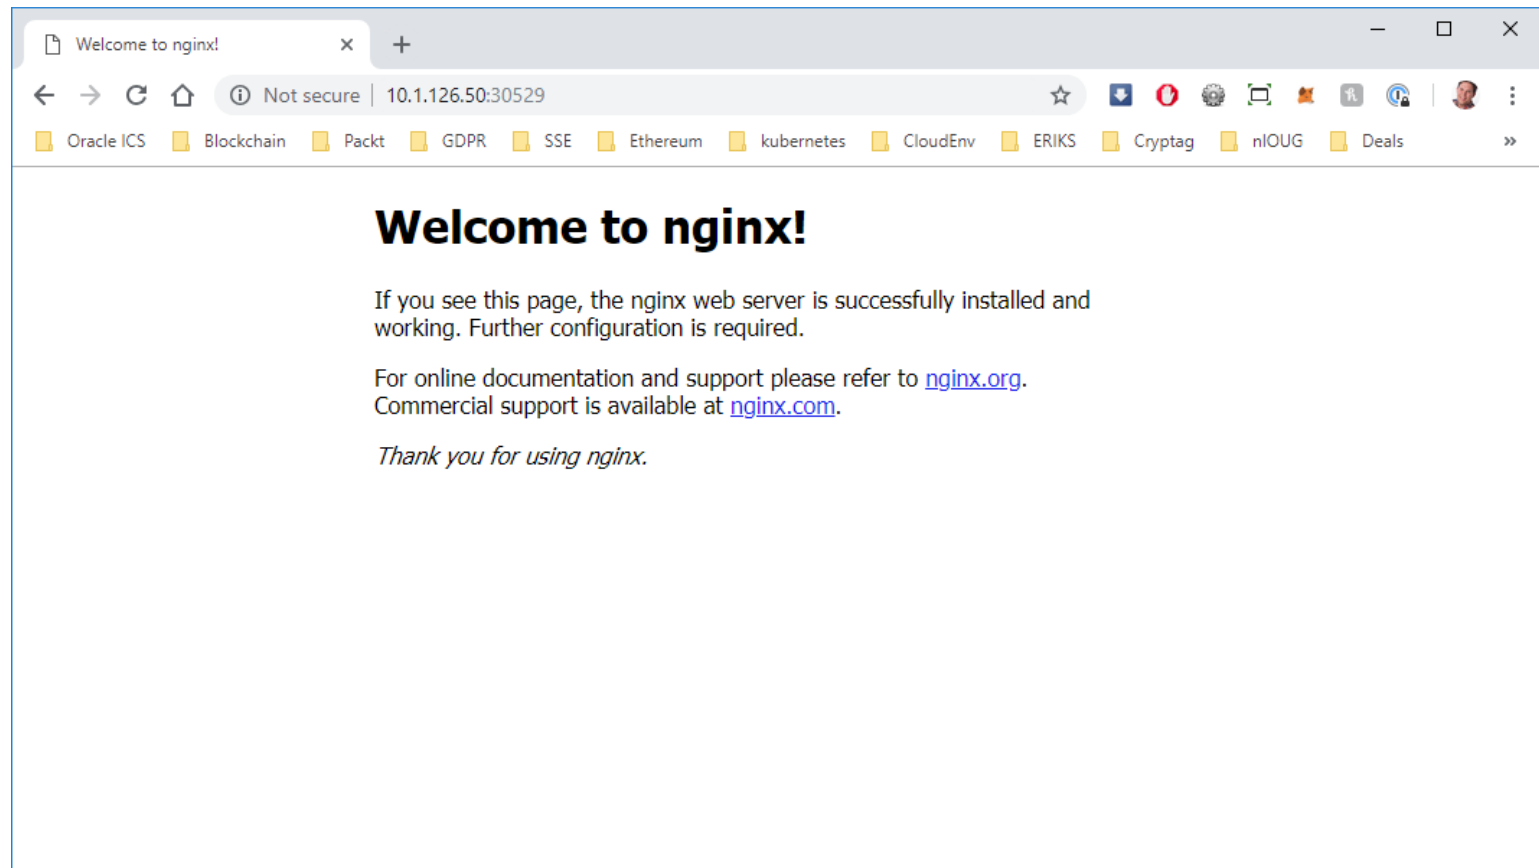
```
minikube service --url=true hello-nginx
```

This command asks `minikube` to return the URL of the `hello-nginx` deployment.

Output

```
http://10.1.126.50:30529
```

4. Go to the url given to see if it is running:



Starting the kubernetes dashboard

Kubernetes has a standard dashboard we can deploy to see what is running on the cluster. When using minikube this dashboard is already available out-of-the-box. Start the dashboard by issuing the following command `minikube dashboard`.


```
$> minikube dashboard
```

```
Enabling dashboard ...
```

```
Verifying dashboard health ...
```

```
Launching proxy ...
```

```
Verifying proxy health ...
```

```
Opening http://127.0.0.1:56817/api/v1/namespaces/kube-system/services/http:kubernetes-dashboard:/proxy/ in  
your default browser...
```

Which if everything goes in order a browser opens the dashboard page, as shown as the image below.

The screenshot displays the Kubernetes Dashboard's 'Overview' page. The top navigation bar includes the Kubernetes logo, a search bar, and a '+ CREATE' button. The left sidebar lists various cluster components: Cluster, Namespaces, Nodes, Persistent Volumes, Roles, Storage Classes, Namespace (set to 'default'), Overview (selected), Workloads, Cron Jobs, Daemon Sets, Deployments, Jobs, Pods, Replica Sets, Replication Controllers, Stateful Sets, Discovery and Load Balancing, Ingresses, Services, and Config and Storage.

The main content area is titled 'Workloads' and features three green circular progress indicators, each labeled '100.00%':

- Deployments**: A table with columns Name, Labels, Pods, Age, and Images. It lists 'hello-nginx' with 1 pod, labeled 'app: hello-nginx', running for 57 minutes using the 'nginx' image.
- Pods**: A table with columns Name, Node, Status, Restarts, and Age. It lists 'hello-nginx-6476b4b4-dnsrp' on the 'minikube' node, in a 'Running' state with 0 restarts, running for 57 minutes.
- Replica Sets**: A table with columns Name, Labels, Pods, Age, and Images. It lists 'hello-nginx-6476b4b4' with 1 pod, labeled 'app: hello-nginx' and 'pod-template-hash: 6476b4b4', running for 57 minutes using the 'nginx' image.

Go and take a look around to see if you can find more details about the hello-nginx deployment/service.

Adding more dashboard functionality

Minikube has a set of built in addons that can be used enabled, disabled, and opened inside of the local k8s environment. Minikube must be running for these commands to take effect.

1. See all addons for minikube using command `minikube addons list`:

```
$> minikube addons list
```

```
- registry: disabled
- registry-creds: disabled
- freshpod: disabled
- addon-manager: enabled
- dashboard: enabled
- heapster: disabled
- efk: disabled
- ingress: disabled
- default-storageclass: enabled
- storage-provisioner: enabled
- storage-provisioner-gluster: disabled
- nvidia-driver-installer: disabled
- nvidia-gpu-device-plugin: disabled
```

One interesting addon is heapster. After enabling heapster you can inspect the health of your environment in much more detail (CPU, Memory, Network usage).

2. Enable heapster addon using command `minikube addons enable heapster`:

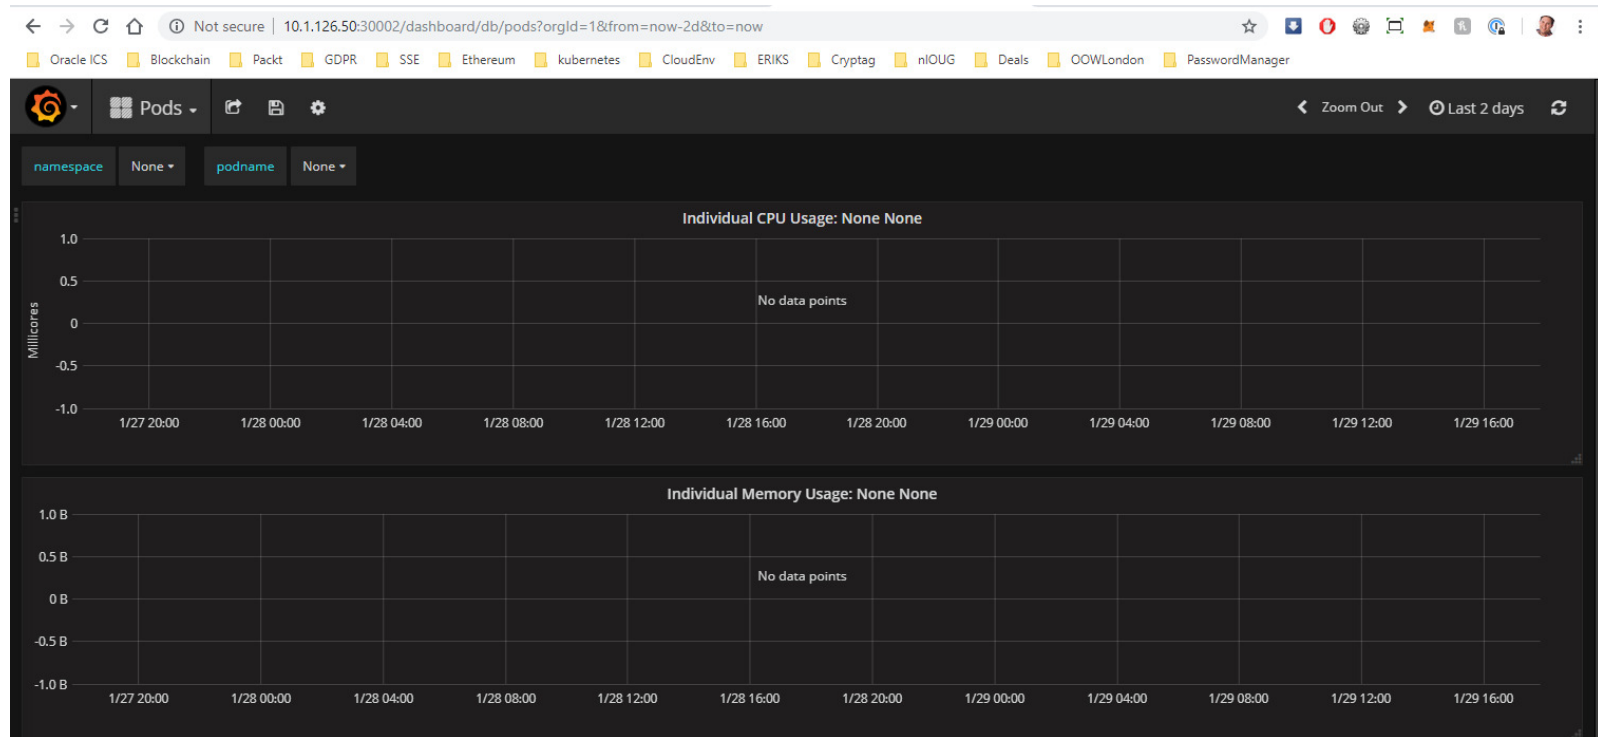
```
$> minikube addons enable heapster
```

```
heapster was successfully enabled
```

3. Open heapster in browser using command `minikube addons open heapster`:

```
# This will open grafana (interacting w/ heapster) in the browser
```

The following image shows graphics collected from pods:



Conclusion

In this lab you got started around with Kubernetes. If you want to play some more you can look at our [Kubernetes Cheat Sheet](#).